

# Operating Systems

## Threads

Slides Courtesy:  
Dr Syed Mansoor Sarwar

# Issues with Processes

- The fork() system call is expensive
- IPC is required to pass information between a parent and its child processes.

# Thread Concept

- A thread is a “lightweight” process which executes within the address space of a process.
- A thread can be scheduled to run on a CPU as an independent unit and terminate.
- Multiple threads can run simultaneously.


# Thread Concept

- Threads have their own
  - Thread ID
  - CPU context (PC, SP, register set, etc.)
  - Stack
  - Priority
  - errno


# Thread Concept

- Threads share
  - Code and data
  - Open files (through the PPFDT)
  - Current working directory
  - User and group IDs
  - Signal setups and handlers
  - PCB

# Single and Multithreaded Processes



single-threaded



multithreaded

# Threads are Similar to Processes

- A thread can be in states similar to a process (new, ready, running, blocked, terminated)
- A thread can create another thread

# Threads are Different from Processes

- Multiple threads can operate within the same address space
- No “automatic” protection mechanism is in place for threads—they are meant to help each other



# Advantages of Threads

- **Responsiveness**
  - Multi-threaded servers (e.g., browsers) can allow interaction with user while a thread is formulating response to a previous user query (e.g., rendering a web page)

# Advantages of Threads

- Resource sharing
  - Process resources (code, data, etc.)
  - OS resources (PCB, PPFDT, etc.)

# Advantages of Threads

- **Economy**

- Take less time to create, schedule, and terminate
- Solaris 2: thread creation is 30 times faster than process creation and thread switching is five times faster than process switching

# Advantages of Threads


- **Performance** in multi-processor and multi-threaded architectures (e.g., Intel's P4 HT)
  - Multiple threads can run simultaneously

# Disadvantages of Threads

- Resource sharing—synchronization needed between threads
- Difficult to write and debug multi-threaded programs

# Single-Threaded Process


```
main()  
{  
    ...  
    f1 (...);  
    ...  
    f2 (...);  
    ...  
}  
  
f1 (...)  
{ ... }  
  
f2 (...)  
{ ... }
```



# Multi-Threaded Process

```
main()  
{  
    ...  
    thread(t1, f1);  
    ...  
    thread(t2, f2);  
    ...  
}  
  
f1(...)  
{ ... }  
  
f2(...)  
{ ... }
```

Process Address Space



# User Threads

- Thread management done by user-level threads libraries
  - Kernel not aware of threads
  - CPU not interrupted during thread switching
  - A system call by a thread blocks the whole process
  - Fair scheduling: P1 has one thread and P2 has 100 threads



# User Threads

- Examples
  - POSIX *Pthreads*
  - Mach *C-threads*
  - Solaris 2 *threads*

# Kernel Threads

- Thread management done by kernel
  - Kernel aware of threads
  - CPU switched during context switching
  - A system call does not block the whole process
  - Fair scheduling: P1 has one thread and P2 has 100

# Kernel Threads

- Examples
  - Windows NT/2000
  - Solaris 2
  - Linux

# Multithreading Models

- Support for both user and kernel threads
- **Many-to-One:** Many user threads per kernel thread; process blocks when a thread makes a system call
- Solaris Green threads
- Pthreads

# Many-to-One Model

User-level  
Threads




Kernel-level  
Thread



# Multithreading Models

- **One-to-One:** One user thread per kernel thread; process does not block when a thread makes a system call
- Overhead for creating a kernel thread per user thread
- True concurrency achieved
- Windows NT/2000, OS/2

# One-to-One Model




# Multithreading Models

- **Many-to-Many:** Multiple user threads multiplexed over a smaller or equal number of kernel threads
- True concurrency not achieved because kernel can only schedule one thread at a time
- Kernel can schedule another thread when a user thread makes a blocking system call
- Solaris 2, HP-UX



# Many-to-Many Model




# **Solaris 2 Threads Model**

- Solaris 2: threads, light-weight processes (LWPs), and processes
  - At least one LWP per process to allow a user thread to talk to a kernel thread
  - User level threads switched and scheduled among LWPs without kernel's knowledge

# Solaris 2 Threads Model

- One kernel thread per LWP; some kernel threads have no LWP (e.g., threads for clock interrupt handler and scheduling)

# Solaris 2 Threads Model



# Pthreads

- A POSIX standard (IEEE 1003.1c) API for thread creation, termination, and synchronization.
- API specifies the behavior of the thread library, implementation is up to developers of the library.
- Common in UNIX operating systems.

# Creating a Thread

- `int pthread_create (pthread_t *threadp,  
                  const pthread_attr_t *attr,  
                  void* (*routine)(void *),  
                  arg *arg);`

# Creating a Thread

## Where:

<code>threadp</code>	The thread we are trying to create—thread ID (TID)
<code>attr</code>	Used to modify the thread attributes (stack size, stack address, detached, joinable, priority, etc.)
<code>routine</code>	The thread function
<code>arg</code>	Any argument we want to pass to the thread function. This does not have to be a simple native type, it can be a 'struct' of whatever we want to pass in.

# Error Handling

- `pthread_create()` fails and returns the corresponding value if any of the following conditions is detected:
  - **EAGAIN** The system-imposed limit on the total number of threads in a process has been exceeded or some system resource has been exceeded (for example, too many LWPs were created).



# Error Handling

- **EINVAL** The value specified by `attr` is invalid.
- **ENOMEM** Not enough memory was available to create the new thread.
- **Error handling:**
  - `#include <errno.h>`
  - Error handling code

# Joining a Thread

- Waiting for a thread
- `int pthread_join(pthread_t aThread, void **statusp);`
- 'statusp' get return value of `pthread_exit`

# Joining a Thread

- Cannot join with a detached thread
- Can only join with thread's in the same process address space
- Multiple threads can join with one thread but only one returns successfully; others return with an error that no thread could be found with the given TID

# Terminating a Thread

- Main thread terminates
- Thread returns
- `void pthread_exit(void *valuep)`
- Returns value pointed to by 'valuep' to a joining thread, provided the exiting thread is not detached

# Example 1

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
/* Prototype for a function to be passed to our thread */
void* MyThreadFunc(void *arg);
int main()
{
    pthread_t aThread;
    /* Create a thread and have it run the MyThreadFunction */
    pthread_create(&aThread, NULL, MyThreadFunc, NULL);
    /* Parent waits for the aThread thread to exit */
    pthread_join(aThread, NULL);
    printf ("Exiting the main function.\n");
    return 0;
}
```

# Example 1

```
void* MyThreadFunc(void* arg)
{
    printf ("Hello, world! ... The threaded version.\n");
    return NULL;
}
```

```
$ gcc hello.c -o hello -lpthread -D_REENTRANT
```

```
$ hello
```

```
Hello, world! ... The threaded version.
```

```
Exiting the main function.
```

```
$
```

# Example 2

```
#include <pthread.h>
```

```
#include <stdio.h>
```

```
#define NUM_THREADS 5
```

```
void *PrintHello(void *threadid)
```

```
{
```

```
    printf("\n%d: Hello World!\n", threadid);
```

```
    pthread_exit(NULL);
```

```
}
```

# Example 2

```
int main (int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc, t;

    for (t=0; t < NUM_THREADS; t++) {
        printf("Creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc) {
            printf("ERROR; return code is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```